

### Assignment -3

#### Python Programming

Assignment Date	19 September 2022
Student Name	Mohana B
Maximum Marks	2 Marks

#### Question-1:

Download the dataset: Dataset

##### Solution:

<https://drive.google.com/file/d/1xkynpL15pt6KT3YSIDimu4A5iRU9qYck/view>

#### Question-2:

Image Augmentation.

##### Solution:

```
from google.colab import drive
drive.mount('/content/drive/')
data_path = '/content/drive/MyDrive/dataset/flowers/'
batch_size = 32
target_size = (64, 64)
```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/content/drive/", force\_remount=True).

```
train_datagen = ImageDataGenerator(rescale=1./255,
                                   shear_range=0.2,
                                   zoom_range=0.2,
                                   width_shift_range=0.1,
                                   height_shift_range=0.1,
                                   horizontal_flip=True,
                                   validation_split=0.2)

test_datagen = ImageDataGenerator(rescale=1. / 255, validation_split=0.2)
```

```
X_train = train_datagen.flow_from_directory(data_path, target_size=target_size, batch_size=batch_size, subset="training", class_mode='categorical')
X_test = test_datagen.flow_from_directory(data_path, target_size=target_size, batch_size=batch_size,
                                         subset="validation",
                                         class_mode='categorical')
```

Found 3457 images belonging to 5 classes.  
Found 860 images belonging to 5 classes.

#### Question-3:

Create Model

##### Solution:

```
model = Sequential()
```

#### Question-4:

Add Layers (Convolution, MaxPooling, Flatten, Dense-(Hidden Layers),Output)

##### Solution:

```

model.add(Convolution2D(32, (3, 3), input_shape=(64, 64, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Convolution2D(32, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Convolution2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())

model.add(Dense(units=64, activation='relu'))
model.add(Dense(units=5, activation='softmax'))
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
conv2d_1 (Conv2D)	(None, 29, 29, 32)	9248
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 12, 12, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 64)	147520
dense_1 (Dense)	(None, 5)	325
=====		
Total params: 176,485		
Trainable params: 176,485		
Non-trainable params: 0		

## Question-5:

### Compile The Model

#### Solution:

```

model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

```

## Question-6:

### Fit The Model

#### Solution:

```
model.fit(X_train, steps_per_epoch=100, epochs=15)
```

```
Epoch 1/15
100/100 [=====] - 29s 271ms/step - loss: 1.3171 - accuracy: 0.4074
Epoch 2/15
100/100 [=====] - 25s 252ms/step - loss: 1.1414 - accuracy: 0.5298
Epoch 3/15
100/100 [=====] - 27s 265ms/step - loss: 1.0823 - accuracy: 0.5658
Epoch 4/15
100/100 [=====] - 26s 257ms/step - loss: 0.9996 - accuracy: 0.5989
Epoch 5/15
100/100 [=====] - 26s 257ms/step - loss: 0.9346 - accuracy: 0.6428
Epoch 6/15
100/100 [=====] - 26s 255ms/step - loss: 0.9086 - accuracy: 0.6485
Epoch 7/15
100/100 [=====] - 26s 255ms/step - loss: 0.8738 - accuracy: 0.6570
Epoch 8/15
100/100 [=====] - 26s 255ms/step - loss: 0.8466 - accuracy: 0.6794
Epoch 9/15
100/100 [=====] - 25s 253ms/step - loss: 0.8086 - accuracy: 0.6904
Epoch 10/15
100/100 [=====] - 25s 251ms/step - loss: 0.7897 - accuracy: 0.6974
Epoch 11/15
100/100 [=====] - 26s 256ms/step - loss: 0.7621 - accuracy: 0.7087
Epoch 12/15
100/100 [=====] - 25s 254ms/step - loss: 0.7526 - accuracy: 0.7043
Epoch 13/15
100/100 [=====] - 26s 257ms/step - loss: 0.7291 - accuracy: 0.7223
Epoch 14/15
100/100 [=====] - 25s 252ms/step - loss: 0.7203 - accuracy: 0.7236
Epoch 15/15
100/100 [=====] - 25s 251ms/step - loss: 0.7175 - accuracy: 0.7274
```

#### Question-7:

#### Save The Model

#### Solution:

```
model.save("model.h5")
```

#### Question-8:

#### Test The Model

#### Solution:

```
def predict():
    img = image.load_img("/content/drive/MyDrive/dataset/flowers/rose/1775233884_12ff5a124f.jpg", target_size=target_size)
    x = image.img_to_array(img)
    x = tf.expand_dims(x,0)

    labels = ['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']

    pred = model.predict(x)
    prediction = labels[np.argmax(pred[0])]

    print(f'The given image is a {prediction}')
    plt.imshow(plt.imread("/content/drive/MyDrive/dataset/flowers/rose/1775233884_12ff5a124f.jpg"))
    plt.axis('off')
    plt.show()
```